

eBay Valuation Bot — decision report

All deep research compiled: Grok + two Claude-chat reports · 17 June 2026

Three calls, made.

- 1. Live data: use eBay's official **Browse API**. It returns current bid, bid count and end time for UK auctions, free up to 5,000 calls/day, no contract. We already built the client.
- 2. What to sell: pivot toward **exact-model cameras and lenses** as the primary lane; keep graded cards as the clean, lower-capital secondary. Avoid watches, sneakers, fashion, Lego, books.
- 3. Business model: **flip yourself**. Don't sell alerts — that market is crowded and has no moat. Your valuation engine plus manual bidding is the edge.

1 · The data problem is solved

Both research passes landed on the same place, and the deeper one verified it field by field. The Browse API's `item_summary/search` returns `currentBidPrice`, `bidCount` and `itemEndDate` for live auctions when you pass `filter=buyingOptions:{AUCTION}`, with `X-EBAY-C-MARKETPLACE-ID: EBAY_GB` for the UK. The old "eBay strips auction data" story is out of date.

Two things matter for us. The Finding API, which old tutorials used for end times, was switched off on 5 Feb 2025 — anything built on it is dead, and Browse is the replacement. And RapidAPI scrapers still can't be verified to return real UK auction timers, which is exactly why the one you bought failed. So that route stays closed.

The lift is small because the code is already written: `--mode browse` calls that endpoint with that filter and header. What's left is access, not coding.

One gap to self-test

Nobody could produce a raw JSON capture of a *live ebay.co.uk auction* with the bid and end-time fields populated. The docs and eBay staff confirm it; a real capture doesn't exist publicly. So the first thing we do once your keyset is live is pull one real GB auction and check those fields come back filled. Quick test, removes the doubt.

2 · What to sell

The verdict: stop treating this as "a card bot." Value on **clean identity**, not on a category. Anything with a machine-readable identity and low condition spread suits the engine; anything where condition is the whole story defeats it.

Cameras and lenses win. Exact model, mount and generation give clean comps; the items are small and dense to store and post; per-flip spreads in recent reseller examples (e.g. £300 in, £550 out) beat the £20–£50 that typical card flips return. UK camera/lens selling fees are also lower (~6.9% vs ~10.9% on collectables).

Graded cards stay as the clean baseline. Slab cert + fixed grade is the easiest thing to value, and postage is trivial. The catch is a softening 2024–25 card market and a crowded field, so the ceiling per flip is lower. Keep it; don't lead with it.

Category	£/flip	Identity	Fit
Cameras & lenses	high	exact model	best
Graded cards / TCG	mid	cert + grade	strong 2nd
Sealed TCG product	thin	SKU	opportunistic
Phones / GPUs	mid	model	scam-prone
Watches	high	subjective	fakes
Sneakers / fashion	mid	condition	avoid first
Lego / books / vinyl	low	mixed	bulk / slow

3 · The bad-listing angle — real, but narrower than the folklore

Your instinct is right, with a sharpening. Underpriced listings do exist and people still profit from them, but the broad "bad photo equals free money" edge has shrunk: buyers and pickers now run Google Lens constantly, and eBay search handles misspellings better than it used to. The edge that survives is specific — a seller who **failed to expose the identity signal that sets the price**: a model suffix buried in a photo, a slab cert left out of the title, a lens mount or version omitted, a misspelled player or set, value hidden inside a vague bundle. That maps onto your tooling far better than messy one-off antiques do, and it's exactly where cameras shine (sloppy listings there usually mean a missing model detail, not total ambiguity). The design consequence is firm: **value automatically only when the listing resolves to a specific identity; everything else goes to manual review, not a confident bid**. We already have that gate — it's the confidence check — so the angle is a tuning job, not a rebuild.

4 · Model assumptions audit — every dial, checked

You asked to look at every assumption in the model. Here's the lot, with what changed. The fixes are now built into the sector profiles in `config.yaml` (run with `--sector`); the base config stays neutral so the tests still hold.

Assumption	Was	Now / verdict
Seller type	private (€0 sell fees)	business — buying to resell with a third-party tool is business activity, and eBay only lets business sellers use such tools. Fixed in every sector.
Final value fee	12.8% flat	category-specific: cards ~10.9%, cameras ~6.9%. Confirm each on eBay's live calculator.
Sold-to-asking ratio	0.85	1.0 — our comps are real sold prices, not asking. Calibrate from outcomes.
USD → GBP rate	0.79, applied to all	cards still USD via Card API (caveat: US sold ≠ UK resale); cameras GBP-only (1.0), no US sourcing.
Comp source	Card API (US, cards only)	fine for cards; no UK sold-comp source wired for cameras yet — the one real open gap.
Min comps gate	8	keep. Matches the high-confidence-or-no-bid rule the research backs.
Max spread gate	0.35	keep; tighten if bad-listing comps come back noisy.
Conservatism k	1.0	keep for cards; consider raising for electronics (returns/scam risk).
Margin floor	25%	keep for cards; for cameras the absolute profit floor matters more.
Profit floor	£15	cameras raised to £40 — high-margin, low-volume means only chasing real spreads.
Price cap	£50	cards £50; cameras £600 (far higher item value).
Title matching	strict tokens + grade regex	works for slabs; cameras need exact-model extraction — to build.
Postage	£3.50 each way	fine for slabs; cameras are heavier — revisit per item.

5 · Where money actually leaks

The research is consistent on how flippers lose. They anchor to active listings instead of sold ones — our engine already values on sold comps, so we're on the right side of that. They buy stock they only think they can move. And on electronics the losses come from returns, item-swaps and false "not as described" claims more than from bad valuation. On cards, the trap is grading overconfidence: raw-to-graded ties up capital for weeks and often doesn't clear. Our model stays read-only and bids to a fee-adjusted max, which sidesteps most of this. The new risk we take on with cameras is condition and returns, so the process needs testing discipline and a conservative bid.

The plan

- 1. **Now:** infrastructure is ready ahead of the key — sector profiles built, business fees fixed, ratio fixed, `--mode browse` waiting. (Done this session.)
- 2. **You:** finish eBay developer sign-up; once approved I set up the account-deletion notification step and wire the keyset.
- 3. **Then:** self-test one live GB auction returns bid + end time; switch the live source to Browse.
- 4. **Build the cameras lane:** exact-model title matching + a UK sold-comp source. Run the cards lane in parallel as the proven baseline.
- 5. **Calibrate:** log real won/resold outcomes, then tune the ratio, floors and fees from data instead of guesses.
- 6. **Later, only if the signal proves out:** consider a small paid alert tier. Not before.

6 · My view on next steps

No new deep-research prompt now. The two reports converge cleanly; another broad sweep would repeat them. The only thing worth a short, targeted look is the one gap above: a reliable **UK sold-price source for cameras**, since the Card API is US/cards-only and eBay's sold data is gated. We can settle that during the build.

On "big mispricing across all sectors" — tempting, but it widens the comp problem instead of narrowing it. Better to nail one clean-identity category first, prove the signal, then clone the profile. Cloning is already a config job.