

eBay Valuation Bot — the pricing model, and where next

Graded sports cards · read-only (it never bids) · 17 June 2026

The bot is now a working tool. You hand it a short watch list of cards; it values each against real sold prices and tells you the most you should pay, how confident it is, and whether the auction you're watching is worth a bid. No live auction feed is needed for this, so it runs today. You still place every bid by hand.

How it prices one card



Start from real sold prices. The Card API returns what graded cards like this one actually sold for, including accepted offers. That beats asking prices, which only say what sellers hoped for.

Take the median, then cut for uncertainty. The median ignores one freak sale. Then the bot subtracts a multiple of the spread, so a scattered set of comps drags the value down on its own. Uncertainty sits inside the number the money decision uses, not in a footnote.

Gate on confidence. Fewer than 8 comps, or too wide a spread, and the bot stays quiet rather than guess.

Work out the ceiling. From the conservative value it takes off your target profit and every cost between winning and reselling: eBay fees, the buyer-protection fee, postage both ways. What's left is the max bid, the most you should ever pay.

Then judge the live price. Under the max with the floors cleared reads **BID**; over it reads SKIP; thin or scattered comps read LOW CONFIDENCE; no comps reads NO DATA.

Worked example. Erling Haaland PSA 10. 11 sold comps, median £95, conservative value £79.05, max bid £49.60. The auction sits at £22, so it's a BID with £27.60 of headroom and high confidence.

Three roads from here

- 1. Sharpen the card model.** Two numbers are still placeholders and probably wrong. The 0.85 ratio was built for asking-price comps, but the Card API already gives sold prices, so it should sit near 1.0. And the 0.79 USD → GBP rate hides a real gap: a US sold price is not a UK resale price. Both need calibrating against your own won-and-resold outcomes. The live auction feed is also still missing, which is why this is manual for now.
- 2. Test a higher-margin, lower-volume category.** Cards were chosen because grading removes condition guesswork, so comps line up cleanly. Electronics can pay far more per flip at lower volume, which fits the effort-per-flip you care about. The risk is that condition variance returns. It only works where identity stays clean: sealed or new stock, exact model numbers, stated refurb grades. Category is already a parameter in the code, so a tight-identity category drops straight in.
- 3. Decide what you're actually selling.** The research consensus is that the edge is sharp buying, not the tech, and that money comes from the flips themselves, from alert tools and Discords, or from writing up what works. The cheapest proven route is to flip with it. Building it into a product only pays if others can't get the same data. Pick one: a flipper who uses a tool, or someone selling the tool.

Code: `targets.py` + `src/valbot/targets.py`, reusing the existing valuation, fee and threshold modules. 33 tests passing. Numbers from the mock fixture; rates confirmed against eBay UK on 16 June 2026. Placeholders (ratio, fx, floors) are starting points to calibrate, not settled values.